

Algorithmic Borders: Human Rights Implications of AI-Driven Visa Adjudication Systems

Ye Min Hong

Columbia University in the City of New York

**“THE TECHNOLOGY IS ALREADY HERE.
THE ACCOUNTABILITY IS NOT.”**

Abstract

The United States has deployed, as of 2026, a suite of artificial intelligence tools across its immigration adjudication system—evidence classifiers, fraud-detection systems, risk-scoring platforms, and translation engines—that screen, classify, and flag visa applications before any human officer reaches a decision. This paper argues that the central failure of this system is not that its automated tools are proven to discriminate, but that they operate without any legally enforceable duty of transparency, audit, or contestability. As a result, algorithmic error and bias can be neither detected nor challenged by the people they affect, and the costs of that opacity fall on the population least equipped to bear them: non-citizens whose only safeguards are administrative, and therefore revocable at a change of administration.

Through a comparative legal analysis of three jurisdictions—the United States, the European Union, and South Korea—the paper shows that the jurisdiction with the most aggressive AI deployment offers the weakest legal protection to those subject to it. It identifies three categories of rights harm implicating the International Covenant on Civil and Political Rights, and argues that the remedy is legislative rather than judicial: the same plenary power invoked to justify the accountability gap is precisely what authorizes Congress to close it. The paper proposes five statutory reforms—mandatory disclosure, explainability, independent auditing, a right to human review, and data-standard requirements—that would bring U.S. practice into alignment with the rights-based frameworks its peers already require.

I. Introduction

On April 8, 2025, I received an email from my university's International Students and Scholars Office informing me that my SEVIS record had been auto-terminated. The stated reason was: "Other — Individual identified in criminal records check and/or has had their visa revoked." The notice did not say how long I had to leave — only that I should depart as soon as possible. No one, including my university's own advisers, could tell me what that meant in practice.

I was not alone. By mid-April 2025, over 1,840 international students and recent graduates from more than 280 colleges and universities had reported SEVIS record shifts — each assigned the same vague "Other" designation, with no citation to a specific provision of the Immigration and Nationality Act, and no explanation of what any of us had done to trigger the action. [1] Universities often discovered that their students' records had been terminated only after checking the SEVIS database themselves. Students, meanwhile, were making consequential decisions — whether to attend class, whether to move out of housing, whether to purchase flights home — with no reliable information about what was actually happening or who had the authority to explain it. That scale of simultaneous, undifferentiated action is what distinguishes this episode from an ordinary administrative error: it is consistent with automated, undifferentiated processing operating at machine speed, without individual review. Whether an automated system in fact drove the terminations has never been disclosed — and that very indistinguishability, from the outside, between algorithmic and human action is itself the first instance of the accountability gap this paper examines: where the two cannot be told apart, the affected person cannot know which occurred, much less contest it.

The case was ultimately resolved through litigation. A temporary restraining order entered on April 18, 2025 by Judge Victoria M. Calvert of the United States District Court for the Northern District of Georgia ordered the government to reinstate affected students' SEVIS records retroactive to March 31, 2025. [2] The government's acknowledgment of error, however, arrived only after the material consequences —

disrupted studies, forfeited employment, and compounded financial costs — had already been absorbed by the students the system was supposed to govern. That sequence — automated action first, legal accountability only later, through costly litigation — is not an anomaly of the 2025 crisis. As Section II demonstrates, it is the operating logic of AI-assisted immigration decision-making more broadly.

What the litigation could not address — and what drew me to the research underlying this paper — was a prior question: how did a system designed to enforce immigration law manage to act on nearly two thousand students simultaneously, without individual review, without legally adequate notice, and without any mechanism for those students to understand, in real time, what had happened to them or what they could do about it? The answer, I came to believe, was not primarily a question of legal authority. It was a question of administrative design. And the design flaw at its center was not unique to the 2025 termination wave — it was structural, systemic, and increasingly technological in character.

As of 2026, USCIS has deployed roughly twenty-nine distinct AI and machine learning systems in its adjudicative workflows, according to the Department of Homeland Security’s publicly available AI Use Case Inventory. [3] The ELIS Evidence Classifier automatically categorizes documents before any human officer reviews a file; USCIS has credited it with saving an estimated twenty-four million page scrolls of reviewer effort over an eight-month period and with nearly doubling the share of cases adjudicated within thirty days, from about 30 to 58 percent. [4] ImmigrationOS — built by Palantir Technologies under a thirty-million-dollar contract with ICE — aggregates data from IRS records, Social Security Administration databases, DMV registrations, passport systems, and license plate readers to construct profiles of immigration status holders and determine enforcement priorities. [5] Fraud detection systems cross-reference submissions against large datasets in ways that practitioners report are generating requests for evidence for documents that were, in fact, submitted. [6] EB-2 National Interest Waiver denial rates rose sharply across FY2025 — reaching 46

percent of adjudicated petitions by the third quarter, up from the mid-30s earlier in the year — though the government has not disclosed what role, if any, automated screening played in that increase. [7]

What remains unpublished — because DHS has not released it — is whether those outcomes distribute evenly across national origin, race, and religion, or whether the algorithms producing them replicate the discriminatory patterns embedded in the historical data on which they were trained. As this paper documents below, that absence of disclosure is not a bureaucratic oversight — it is the central legal and human rights problem this paper addresses.

This paper argues that the central failure of U.S. immigration adjudication is not that its automated systems are proven to discriminate, but that they operate without any legally enforceable duty of transparency, audit, or contestability — so that algorithmic error and bias can be neither detected nor challenged by the people they affect. Because the only safeguards that have ever applied to these systems are administrative, and therefore revocable at a change of administration, the remedy this paper defends is not judicial but legislative: the same plenary power invoked to justify the gap is precisely what authorizes Congress to close it, through the transparency, audit, and human-review obligations that the European Union and South Korea already require.

The argument proceeds in four sections. Section II maps the operational landscape of AI-driven visa adjudication, examining the key systems currently deployed, their scale, and their documented effects on outcomes. Section III conducts a comparative legal analysis across three jurisdictions — the United States, the European Union, and South Korea — assessing the adequacy of each framework to protect individuals subjected to AI-assisted adjudication. Section IV identifies three specific categories of rights harm: the unauditable risk of algorithmic bias under ICCPR Article 26, the transparency deficit and its relationship to due process, and the erosion of meaningful procedural

rights through the “decision support” framing. Section V proposes five legislative reforms that would bring U.S. practice into alignment with the rights-based frameworks developed by peer jurisdictions.

The claim of this paper is not that AI has no place in immigration administration. The scale of the U.S. system — over eleven million pending cases as of mid-2025 [8] — makes some degree of automated assistance operationally defensible. The claim, rather, is that the deployment of such tools without transparency, accountability, or meaningful recourse for those they affect is not an administrative oversight. It is a structural choice that transfers the full cost of algorithmic error onto the population least equipped to absorb it: the visa applicants, students, and immigrants whose legal status depends on systems they cannot see, cannot challenge, and are never told exist. I know what that cost looks like from the inside. This paper is an attempt to trace what it looks like from the outside — and to argue that the law should do better.

II. The Operational Landscape of AI-Driven Visa Adjudication

A. From Bureaucratic Processing to Algorithmic Triage

The United States Citizenship and Immigration Services has undergone a fundamental operational shift. What was once a paper-intensive, officer-driven review process has increasingly become an algorithmically mediated one, in which artificial intelligence tools screen, categorize, and flag cases — in a growing share of the workflow, before any human adjudicator reaches a substantive decision. The Department of Homeland Security publishes an AI Use Case Inventory listing every artificial intelligence and machine learning system deployed across DHS agencies; the January 28, 2026 update includes roughly twenty-nine distinct USCIS AI use cases — a figure drawn directly from the government’s own public disclosure, not from practitioner estimates. [9] These are not pilot programs or future-facing initiatives — they are embedded in daily adjudicative workflows. This

paper centers on that adjudicative layer — the systems that screen, classify, and score benefit applications submitted by people seeking or holding lawful status. It treats the enforcement tools described below, such as ImmigrationOS and ICE’s risk-scoring systems, not as its primary subject but as the more extreme expression of the same accountability gap: the legal posture of a person already in removal proceedings differs from that of a visa applicant, yet both are governed by automated processing that operates without any duty to disclose, explain, or justify what it has done.

A note on the evidentiary basis of what follows is warranted, because the same opacity this paper criticizes also constrains the evidence available to study it. Government disclosures establish that these systems exist, at what scale, and — in the case of the ELIS Evidence Classifier — with what efficiency gains. What those disclosures do not provide is per-case performance data: error rates, the weight automated outputs carry in individual adjudications, or their distributional effects across national origin, race, and religion. For the operational effects described below, this paper therefore relies in part on practitioner reporting — accounts by immigration attorneys and legal-services organizations of what they observe in their cases — rather than on primary agency data. Such reporting is probative but not conclusive; it is flagged as practitioner observation throughout, and the resulting account of how these systems affect applicants should be read as indicative of documented mechanisms rather than as a quantified measure of their aggregate impact. That evidentiary gap is not incidental to the argument. It is a direct consequence of the disclosure vacuum this paper identifies as the central problem, and it is precisely why the reforms proposed in Section V center on the auditing and transparency obligations that would make a more definitive account possible.

DHS maintains that human adjudicators retain final decision-making authority, and that AI systems do not grant or deny immigration benefits. [10] That formal position, however, obscures a more complex reality: automated tools shape which evidence surfaces for review,

which inconsistencies get flagged, and which cases arrive at a human officer's desk already marked for heightened scrutiny. Final authority and sole influence are not the same thing.

B. Key AI Systems Currently Deployed

1. ELIS Evidence Classifier

The Electronic Information System (ELIS) Evidence Classifier is a machine learning tool that automatically categorizes and indexes documents submitted with immigration petitions. Over the period from September 2021 to May 2022, USCIS credited the tool with saving roughly twenty-four million page scrolls of reviewer effort — on the order of 13,000 hours — and with nearly doubling the share of cases adjudicated within thirty days, from about 30 to 58 percent. These are efficiency gains, not measures of accuracy: USCIS has not published any error-rate data, and practitioners report requests for evidence for documents that were in fact submitted, consistent with classifier mis-tagging. [11]

Per the 2024 USCIS Privacy Impact Assessment, the Evidence Classifier categorizes and surfaces documents for human review; it does not make adjudication decisions. Practitioners describe a related downstream effect: when a document bundle is mis-scanned, out-of-order, or illegible, the classifier's confidence in its categorization drops, and the adjudicating officer is correspondingly more likely to request a secondary evidence submission. [12] This means a purely technical failure in the AI layer can generate a procedural burden on the applicant that mirrors the effect of a substantive evidentiary deficiency.

2. ImmigrationOS

ImmigrationOS represents a qualitatively different category of AI intervention. In April 2025, ICE awarded Palantir a thirty-million-dollar contract to build ImmigrationOS, a unified case-management and analytics platform that aggregates IRS tax records, Social Security Administration data, DMV registrations, passport and visa data, and

license plate reader data. [13] Federal officials have said the platform draws these disparate records together to prioritize enforcement targets and to track individuals as they move through the immigration system. [14]

Prior visa applications, travel histories, and business registrations may be cross-referenced in ways that were not previously routine. Errors that once went unnoticed may now trigger delays, requests for evidence, or referrals for further review. [15]

3. StateChat and the Document Translation Service

StateChat, a generative AI chatbot the State Department launched in late 2024, is used by consular and departmental staff to interpret internal guidance, draft communications, and coordinate across posts — functions that increasingly touch the visa-issuance process. [16]

Separately, USCIS officers can now upload foreign-language documents and receive near-instant AI-generated translations via Microsoft Azure AI Translator. Whether that translation is preserved in the administrative record, and in what form, remains an open question — one that matters most in proceedings where credibility may hinge on precise wording. [17]

4. Fraud Detection Systems (FDNS) and Babel X

USCIS and its Fraud Detection and National Security Directorate rely on AI systems to flag anomalies and cross-reference filings, while Customs and Border Protection uses Babel X, a tool that mines social media and other publicly available data to build profiles of travelers and screen them at ports of entry. [18]

C. Measurable Consequences for Applicants

The cumulative effect of these systems on adjudication outcomes is real, but the denial-rate trend and the role of automation within it must be kept analytically distinct. The trend itself is now a matter of public record: USCIS Form I-140 data show the EB-2 NIW denial rate climbing across FY2025, from roughly 37 percent of adjudicated petitions in the first quarter to 46 percent in the third; for the fourth quarter of FY2025, a National Foundation for American Policy analysis of subsequent USCIS

data places the rate at 64.3 percent, with EB-1A denials climbing in parallel. Analysts and a former USCIS adjudicator attribute this increase primarily to a more exacting application of the Dhanasar standard, rather than to any announced change in the governing legal test or to automation as such. What cannot be established from the public record is the contribution of AI-assisted screening to that trend, because USCIS has released no data isolating it. The role of automation is documented instead at the level of mechanism: practitioners report that mis-tagged evidence generates requests for evidence for documents that were in fact submitted, and that cross-document mismatch detection causes small discrepancies in job titles, dates, or signatures to trigger automated flags. Whether these operational effects move the aggregate denial rate, and by how much, is precisely the question that the absence of disclosure makes impossible to answer. [19]

Automation does not need to ‘decide’ the case to materially affect it. [20] This observation captures the central problem this paper examines: a system in which no single automated decision is formally determinative, but in which algorithmic processing at every stage of the review workflow can systematically disadvantage applicants who lack knowledge of, access to, or recourse against the processes shaping their cases. The question that follows is whether existing legal frameworks are equipped to govern a system that operates this way — and, as Section III demonstrates, the answer across all three jurisdictions examined is inadequate, and in the United States, critically so.

III. A Comparative Legal Analysis of AI Governance Frameworks

The legal landscape governing AI-assisted adjudication varies considerably across jurisdictions. An examination of three frameworks — the United States, the European Union, and South Korea — reveals not only divergent regulatory approaches but also a fundamental asymmetry: the jurisdiction with the most aggressive AI deployment in immigration adjudication has the weakest legal protections for those affected by it.

A. The United States: Broad Deference, Narrow Protections

1. The Plenary Power Doctrine and Its Limits

American immigration law rests on a doctrine that has no parallel in domestic civil law: the plenary power doctrine, which holds that Congress exercises nearly unreviewable authority over immigration matters. Long-standing Supreme Court precedent has recognized Congress as having plenary power over immigration, giving it almost complete authority to decide whether foreign nationals may enter or remain in the United States. [21] This structural deference to the political branches substantially constrains the judiciary's ability to scrutinize the procedural mechanisms through which immigration decisions are made.

The doctrine does not, however, extinguish due process protections entirely. The Supreme Court has held that aliens physically present in the United States, regardless of their legal status, are recognized as “persons” guaranteed due process of law by the Fifth and Fourteenth Amendments, and that even those whose presence is unlawful, involuntary, or transitory are entitled to that constitutional protection. [22] The critical word, however, is “present.” The due process protections that attach to persons within U.S. territory thin considerably at the border and dissolve almost entirely for visa applicants abroad.

2. The Administrative Procedure Act and Its Gaps

The Administrative Procedure Act (APA) provides a second avenue for challenging agency action, requiring that final agency decisions be supported by reasoned explanation and be neither arbitrary nor capricious. Where AI tools contribute to a denial, however, the APA's utility is limited: AI-flagged cases may move through the adjudicative pipeline in ways where automation influences the front end of adjudication — shaping timelines, scrutiny levels, and documentation burdens — without constituting a “final agency action” subject to APA review. [23]

3. Assessment: A Protective Gap

The United States has no federal statute requiring transparency, auditing, or explainability for AI systems used in immigration adjudication. The Court remains implicitly willing to give the political branches leeway to use immigration authorities in rough-hewn ways. And lower courts that might otherwise expand individual-rights review are constrained by the expectation that they would overvalue individual interests and undervalue governmental ones. [24] In the absence of legislative action, applicants affected by AI-assisted adjudication have no right to know that automated tools were involved, no right to receive an explanation of how their file was processed, and no meaningful avenue for challenging algorithmic errors that shaped — but did not formally decide — their cases. The only accountability framework that has ever applied to immigration AI has been administrative rather than statutory — a sequence of OMB memoranda that a single change of administration revised and, in 2025, replaced (see Section IV.B.2, *infra*). A safeguard that one executive can install and the next can rescind is not, in any meaningful sense, a legally enforceable protection for the applicants it nominally covers.

B. The European Union: A Rights-First Architecture

1. GDPR Article 22: The Right Against Automated Decision-Making

The General Data Protection Regulation, in force since 2018, establishes in Article 22 a right not to be subject to decisions based solely on automated processing that produce legal or similarly significant effects. Article 22 restricts solely automated decisions with legal or similarly significant effects unless strict conditions and safeguards are met, typically requiring meaningful human involvement. [25]

The practical significance of this protection has been expanded by EU judicial interpretation. The European Court of Justice's *Schufa* decision reinforced a broad reading of Article 22, holding that generating a score through automated profiling can itself constitute an automated decision where third parties rely heavily on that score —

meaning that AI-generated rankings or flags that effectively determine AI-generated rankings or flags that effectively determine outcomes may be treated as automated decisions even if a human nominally approves the result. [26]

2. The EU AI Act: High-Risk Classification and Structural Obligations

The EU AI Act, which entered into force in August 2024, creates a second layer of protection. Its high-risk regime — including the Article 86 right to explanation — was originally scheduled to apply from August 2026 but was deferred to December 2027 for standalone Annex III systems, which include border-control and law-enforcement decision-support tools, under the Digital Omnibus, Regulation (EU) 2026/1744, a delay that itself illustrates how demanding these obligations are to operationalize. Article 86 introduces an individual right: any person subject to a decision significantly affected by a high-risk AI system is entitled to a clear explanation of the AI's role in that decision, the main parameters that influenced the output, and the level of human oversight involved — a right that extends beyond GDPR Article 22, which covers only fully automated decisions. [27]

3. Assessment: A Model With Structural Limitations

The EU framework represents the most developed effort to subject AI systems to rights-based constraints. Its primary limitation, for purposes of this paper, is geographic: its protections extend to individuals within the EU, and visa applicants from outside that jurisdiction may not benefit from its provisions to the same degree. Moreover, the AI Act and GDPR are complementary instruments that must both be satisfied: compliance with one does not discharge obligations under the other, and failure to comply with either can result in separate penalties. [28]

C. South Korea: A Progressive Framework With Unresolved Gaps

1. The AI Framework Act: Structure and Scope

South Korea enacted the Framework Act on Artificial Intelligence Development and Establishment of a Foundation for Trustworthiness on

January 21, 2025, with full enforcement beginning January 22, 2026. Although the EU was first to pass comprehensive AI legislation, South Korea became the first country to bring such legislation into force, as the EU's high-risk obligations — now deferred to December 2027 — remain pending. [29]

The Act's regulatory architecture centers on the concept of "high-impact AI," defined as systems that may significantly affect or pose risks to human life, physical safety, and fundamental rights, and that operate in sectors including healthcare, energy, transportation, and public services. The Act stipulates duties to ensure AI transparency and AI safety, and sets forth specific obligations for operators of high-impact AI systems. [30]

2. Core Obligations: Articles 31, 32, and 34

Under Article 31, AI business operators offering high-impact or generative AI products must notify users in advance that the product or service utilizes AI. [31] Article 32 requires operators of AI systems exceeding prescribed computational thresholds to identify, assess, and mitigate risks throughout the AI lifecycle. Operators involved with high-impact AI are obligated to establish and implement risk management measures and user protection plans, and to ensure human oversight and supervision of high-impact AI systems under Article 34. [32]

3. Assessment: An Incomplete Application to Immigration

South Korea's AI Framework Act constitutes a meaningful advance over the U.S. approach. For a Korean visa applicant processed by a USCIS AI system, however, the Act provides no direct protection: it governs operators within South Korean jurisdiction, not U.S. federal agencies. The Act has been criticized for prioritizing industrial promotion over proactive regulation, and for lacking adequate protective measures compared to the EU AI Act, which imposes stricter sanctions on AI technologies that may be ethically or socially harmful. [33]

D. Comparative Summary

The three frameworks reveal a pattern that is troubling rather than reassuring. The jurisdiction with the most extensive AI deployment in immigration adjudication — the United States — has the least developed legal architecture for protecting applicants' rights within that system. The EU has constructed the most robust rights-based framework, but its geographic scope limits its relevance to applicants processed by U.S. agencies. South Korea has demonstrated political will to regulate AI comprehensively, but has not yet extended those protections explicitly to the immigration context or to its own nationals processed abroad. The gap that emerges from this comparative analysis is not merely a technical regulatory deficit. It is a structural asymmetry in which the people most affected by AI-assisted adjudication — visa applicants, asylum seekers, and migrants — are also the least protected by the legal frameworks currently governing that adjudication.

To see what this asymmetry means in practice, consider a hypothetical applicant — a petitioner for an EB-2 National Interest Waiver whose supporting documentation is assigned a low-confidence tag by an automated classifier of the kind described in Section II, prompting a request for additional evidence. The automated input is the same in each jurisdiction; only the legal consequences differ. In the European Union, she could invoke Article 22 of the GDPR — in force since 2018 — to demand an explanation of the automated role in her case, and, once the AI Act's Article 86 takes effect, a further right to explanation specific to high-risk systems. In South Korea, the Framework Act's notice and human-review obligations would entitle her to be told that an automated system had been used and to seek human reconsideration. In the United States, she would be entitled to neither: no statute requires the agency to disclose that automation played any part, and no mechanism exists for her to contest it. The same fact pattern, run

through three legal systems, yields three different levels of protection — and the jurisdiction subjecting her to the most automated processing offers her the least recourse.

IV. Human Rights Implications of AI-Driven Visa Adjudication

The structural asymmetry identified in Section III — the most AI-intensive jurisdiction offering the least legal protection — has concrete human rights consequences. This section identifies three specific categories of rights harm that flow from the regulatory vacuum documented above: the unauditable risk of algorithmic bias and the right to non-discrimination under ICCPR Article 26 (Part A); the transparency deficit and its relationship to due process (Part B); and the erosion of meaningful procedural rights through the “decision support” framing (Part C). Taken together, these three categories demonstrate that the costs of the current regulatory gap are not abstract — they fall on identifiable people, in identifiable ways, with legal consequences that existing frameworks have failed to address.

A. The Unauditable Risk of Algorithmic Bias and the Right to Non-Discrimination

1. The Technical Mechanism of Bias

Algorithmic bias is not a hypothetical risk; it is a well-documented failure mode of systems trained on historically generated data. Whether the specific tools deployed in U.S. immigration adjudication exhibit it cannot be established on the public record — not because the risk is absent, but because DHS has disclosed no data that would allow it to be detected or audited. Discrimination may occur at any point in the AI lifecycle as a consequence of data, social, or technical bias, or the mode of implementation. When training data are not representative of the population under consideration, or when they embed real-world discrimination, such discrimination is transmitted to the AI system and

disproportionately impacts those already experiencing historical disadvantage. [34]

Commentators warn that AI migration tools, trained on data that encodes historical discrimination and built by developers without expertise in racial bias, operate within legal systems that offer little transparency or recourse — and so risk reproducing the very inequities they claim to eliminate. On this view, absent robust safeguards, such systems can entrench racialized exclusion while presenting themselves as neutral efficiency measures. [35]

2. Concrete Examples: Risk Scoring Without Accountability

In 2024, more than 140 advocacy groups asked DHS to suspend select immigration AI pilots, arguing that ensuring fair, unbiased decision-making is difficult without full transparency into the algorithms that power tools like the ICE Hurricane Score and Risk Classification Assessment (RCA). [36] DHS characterizes these tools as informing rather than making custody and supervision decisions; critics counter that in practice they shape whether migrants are released or subjected to electronic monitoring. What is not disputed is that the algorithm's role is disclosed neither to the affected individual nor to their legal counsel. [37]

3. The International Human Rights Framework: ICCPR Article 26

The applicable international standard is Article 26 of the International Covenant on Civil and Political Rights (ICCPR), which guarantees all persons equal and effective protection against discrimination on any ground including race, national origin, and religion. The United States has ratified the ICCPR, accepting obligations that extend beyond its domestic constitutional guarantees. Those obligations are not, on their own, enforceable by an individual in a U.S. court: the Senate ratified the Covenant subject to a declaration that its substantive articles are non-self-executing. That limitation, however, cuts in favor of the argument advanced here rather than against it. Because the ICCPR binds the United States as a matter of state responsibility while

conferring no privately enforceable right, it is precisely Congress — not the judiciary — that must give its guarantees domestic effect. The Covenant thus does not supply a cause of action for the visa applicant; it supplies a measure of what the legislature owes her. [38]

What makes regulation of algorithms particularly complicated is the problem of procedural invisibility: algorithms arrive at outcomes through processes that are almost impossible to reconstruct, even for their designers. The input of data may not be discriminatory on its face, but may reflect existing patterns of structural discrimination — a complexity that cannot be adequately addressed by the broad, general definitions of non-discrimination as they appear in ICCPR and other international human rights instruments. [39]

B. The Transparency Deficit and the Right to Know

1. Algorithmic Opacity in Legal Context

Transparency is not merely a technical aspiration in the context of government decision-making — it is a legal requirement grounded in due process, reasoned decision-making, and the right to contest adverse determinations. Procedural justice requires that affected parties understand the grounds for decisions affecting them and have meaningful opportunity to contest adverse determinations. The inadequacy of existing disclosure mechanisms has been documented in ongoing federal litigation: advocacy organizations have sued DHS and USCIS under the Freedom of Information Act specifically to compel disclosure of the agency's AI use policies and systems. [40]

A well-known problem with AI systems is their non-explainability: the reasoning underlying a given output may be inaccessible even to the system's developers. And even if developers and users understand what is behind a decision, that understanding may not be transmissible to the affected traveler in any meaningful form. Moreover, automated risk assessment systems can and do make mistakes. [41]

The legal significance of this opacity was recognized in U.S. case law: a

court found that an algorithmic scoring system used to rank teachers had to be open to scrutiny in order to permit affected teachers to verify the accuracy of their score and challenge the decision if they disagreed. Without access to the underlying equations, codes, decision rules, and assumptions, those affected could not exercise their constitutionally protected due process rights — their scores remained a “black box, impervious to challenge.” [42]

2. The Absence of a Federal Explainability Obligation

The Trump administration’s January 2025 Executive Order rescinding the Biden administration’s AI governance framework eliminated the primary federal mechanism for requiring transparency in government AI systems. In its place, OMB Memorandum M-25-21 (April 2025) rescinded and replaced M-24-10, folding the earlier “rights-impacting” and “safety-impacting” categories — and the individualized review they triggered — into a single “high-impact AI” definition oriented toward accelerating adoption rather than constraining it. That successor framework still treats AI used to assess a foreign national’s entry into or continued presence in the United States as presumptively high-impact, and it retains an instruction that agencies offer consistent remedies or appeals. But the memorandum provides on its face that it creates no enforceable right or benefit and governs only agencies’ own use of AI — so neither the high-impact designation nor the remedies instruction gives a visa applicant anything to invoke. [43]

There is no federal statute, no applicable executive order, and no agency regulation requiring USCIS or ICE to disclose to visa applicants that their files have been processed by automated tools, what those tools concluded, or what weight that automated processing was given by the human officer who ultimately signed the decision. DHS has not published outcome-based performance metrics comparing AI-assisted processing to traditional processing, nor has it released data evaluating whether automated assistance changes approval or denial rates. [44]

3. Contrast With the EU Standard

The gap becomes most visible in comparative perspective. The EU

affords such applicants a right to explanation that the U.S. framework lacks: the GDPR's Article 22 already constrains solely automated decisions, and the AI Act's Article 86 will, once in force, add a right for any person subject to a significantly affecting high-risk decision to a clear explanation of the AI's role in the decision, the main parameters that influenced the output, and the level of human oversight involved. [45] The U.S. framework offers applicants no equivalent right.

C. Due Process Erosion and the Right to a Fair Hearing

1. The Procedural Stakes

The right to a fair hearing before an adverse government determination is a core principle of both U.S. constitutional law and international human rights law. ICCPR Article 14 guarantees the right to a fair and public hearing by a competent, independent, and impartial tribunal. AI-assisted adjudication undermines these guarantees in ways that are not fully captured by the formal preservation of human decision-making authority. When AI tools pre-sort, pre-classify, and pre-flag applications, the human officer who reviews the resulting file exercises judgment within a decision space that has already been shaped by automated processing that the applicant had no opportunity to contest, correct, or even observe.

2. The “Decision Support” Fiction

The “decision support” framing — under which AI tools influence but do not formally decide cases — obscures the due process problem rather than resolving it. It keeps the legal accountability framework anchored to a human decision-maker, even as the substantive factors shaping that decision are generated by systems operating outside any procedural safeguard. The ELIS Evidence Classifier, described in Section II, shows how this works in practice. Consider an EB-2 petition supported by documentary evidence. When the classifier assigns a low-confidence tag to one of those documents — because the scan is poor or the pages are out of order — the officer who inherits the file becomes measurably more likely to issue a request for evidence, and the petitioner must then rebut a deficiency that originates not in the merits of her case but in an

automated mis-read she never sees. No official has formally “decided” anything against her, and no adverse determination appears in the record; yet her evidentiary burden, her timeline, and her odds of success have all shifted. That is why the “decision support” label operates as a fiction: it locates accountability in a human choice that has already been pre-shaped by an unseen process, and it is this structural opacity that threatens the right to effective judicial protection. The same pattern recurs in enforcement. When an automated risk assessment flags a traveler, officers may subject them to more extensive inspection, and the traveler has no practical means of ascertaining why that flag was generated. [46]

3. The Compounding Effect on Vulnerable Populations

These due process deficits do not distribute evenly. They fall most heavily on applicants who lack legal representation, who submit from countries with limited access to qualified immigration counsel, and who do not read English fluently enough to parse the technical requirements of a request for evidence or notice of intent to deny. Practitioners report that AI-driven data-matching systems can produce false mismatches, leading to rejections of properly filed petitions; that mis-tagged evidence can generate requests for evidence for documents that were actually submitted; and that cross-document mismatch detection can cause small discrepancies in job titles, dates, or signatures to trigger automated flags. [47] For applicants who cannot identify the source of a deficiency notice, the gap between formal procedural rights and practical access to those rights widens into something close to a denial of meaningful process.

V. Conclusion and Policy Recommendations

A. Principal Findings

This paper has traced a specific and consequential arc. The United States government has deployed, as of 2026, a suite of AI tools — ELIS Evidence Classifier, ImmigrationOS, StateChat, Babel X, and various

fraud detection and risk scoring systems — that collectively reshape the processing of visa applications before any human adjudicator exercises meaningful independent judgment. These tools influence which documents surface for review, which applicants are flagged for heightened scrutiny, and which files arrive at a human officer’s desk already pre-sorted by algorithmic criteria that the applicant cannot observe, contest, or correct.

The comparative legal analysis in Section III demonstrated that this deployment is occurring in a regulatory environment that is structurally the weakest of the three jurisdictions examined. The European Union has constructed a layered framework — GDPR Article 22 and the EU AI Act — that provides affected individuals with substantive rights to explanation, human review, and contestability of AI-influenced decisions. South Korea’s AI Framework Act, in force as of January 22, 2026, establishes mandatory transparency obligations, safety requirements, and human oversight duties for high-impact AI systems. The United States, by contrast, has no federal statute requiring disclosure, explainability, auditing, or human review rights in the context of AI-assisted immigration adjudication.

The human rights analysis in Section IV identified three specific categories of rights harm: the unauditable risk of algorithmic bias, which where present would violate ICCPR Article 26’s guarantee of equal protection; a transparency deficit that undermines both the procedural right to a reasoned decision and the substantive ability to contest adverse outcomes; and the erosion of meaningful due process through the “decision support” framing that preserves formal human authority while allowing algorithmic processing to shape decision outcomes outside any procedural safeguard.

B. Policy Recommendations

The following five recommendations address, in turn, the three categories of harm identified in Section IV. Recommendations 1 and 2 respond to the transparency deficit; Recommendation 3 addresses

algorithmic bias; Recommendation 4 responds to due process erosion; and Recommendation 5 addresses the data conditions that generate both bias and opacity. Each is drawn from existing legislative proposals or comparative international models and is designed to be actionable within the current U.S. statutory framework.

Recommendation 1: Mandatory Disclosure and Notice

Congress should enact legislation requiring that any applicant whose immigration file has been processed, classified, flagged, or scored by an automated system receive written notice of that fact as part of any request for evidence, notice of intent to deny, or denial. The notice should identify the AI tools involved, describe in plain language the type of processing performed, and explain what role, if any, that processing played in the adjudicating officer's review. This recommendation draws on the model established by the White House Blueprint for an AI Bill of Rights, which advanced requirements for notice to individuals and explanations of automated decisions. [48]

Recommendation 2: Statutory Explainability Requirements

Congress should require that USCIS and ICE maintain documentation of the logic, parameters, and training data characteristics of all AI systems used in immigration adjudication, and make this documentation available for inspection in proceedings where an applicant challenges an adverse determination. A U.S. court has already recognized that where an algorithmic scoring system affects individuals' rights, those individuals must have access to the underlying equations, decision rules, and assumptions in order to exercise their constitutionally protected rights — without which the AI system remains a mechanism of adjudication that is, in practice, impervious to challenge. [49]

Recommendation 3: Independent Algorithmic Auditing

Congress should establish a mandatory, periodic, independent audit regime for AI systems used in immigration adjudication. Many state AI bills and proposed federal legislation impose affirmative reporting requirements and bias audit mandates, requiring impact assessments

to check for discriminatory effects against protected groups. [50] A federal immigration-specific audit mandate should require auditors — drawn from a roster of independent technical and civil rights experts, not agency personnel — to assess whether AI tools produce statistically differential outcomes across national origin, race, religion, or other protected characteristics, and to publish those findings publicly.

Recommendation 4: A Right to Human Review

Congress should codify a right for visa applicants to request de novo human review of any determination that was preceded by an automated flag, risk score, or AI-generated evidentiary classification. This right should be available without cost to the applicant, should suspend any adverse determination pending completion of the review, and should require the reviewing officer to document independently, without reference to prior AI outputs, the evidentiary basis for their decision. GDPR Article 22 already establishes this principle in EU law, requiring meaningful human involvement as a safeguard against decisions based solely on automated processing. [51]

Recommendation 5: Data Minimization and Bias Mitigation Standards

Congress should require that AI systems deployed in immigration adjudication be trained only on data that has been reviewed and certified for demographic representativeness, and that DHS publish, on an ongoing basis, outcome data disaggregated by national origin, race, and religion for each AI tool in use. The Equal Employment Opportunity Commission has recognized that if an algorithm results in a disparate impact on protected classes, liability attaches regardless of whether the tool was internally developed or procured from a third-party vendor. [52] The same principle should apply to government AI tools in immigration adjudication.

The strongest objection to this paper's recommendations is structural rather than empirical. One might argue that the current state of American immigration AI reflects not an oversight but a considered constitutional design: that Congress exercises plenary power over

immigration — nearly unreviewable by the courts — precisely because immigration involves sovereign judgments about national membership that the political branches are best equipped to make. On this view, the roughly eleven-million-case USCIS adjudication backlog [8] makes some degree of automated processing a practical necessity, and the absence of judicially enforceable transparency obligations reflects a deliberate legislative choice to leave the executive branch free to modernize adjudication without judicial second-guessing. Fraud detection and national security screening, the argument continues, depend on confidentiality that disclosure mandates would compromise.

This objection, pressed to its conclusion, refutes itself. The plenary power doctrine constrains judicial review — it does not constrain Congress. The recommendations offered in this paper are directed not at courts but at the legislature: they ask Congress to exercise the very authority that the plenary power doctrine confirms it possesses. A statutory transparency mandate, an auditing requirement, or a data-disclosure obligation for immigration AI is not a judicial incursion into executive prerogative; it is an act of Congress regulating the administrative apparatus that Congress itself created. Far from being blocked by plenary power, legislative reform of the kind proposed here is precisely what the doctrine contemplates: political-branch accountability for political-branch decisions. The objection thus becomes a reason to pass legislation, not a reason to forgo it.

Nor does the confidentiality interest in fraud detection and national security screening argue against accountability as such: transparency can be designed with security in mind — disclosure to cleared counsel under protective order, in camera review, and audit findings reported in aggregate rather than case by case — so that the choice is not between secrecy and oversight, but between calibrated disclosure and none. And while the domestic due process protections invoked here reach applicants within U.S. territory but thin for those abroad, the non-discrimination and fair-hearing obligations the United States accepted in ratifying the ICCPR are not territorially bounded in the same

way — so the largest affected class is not, as the objection assumes, beyond the reach of any enforceable norm. And the efficiency argument, however weighty, proves only that AI processing should be permitted — not that it should be unaccountable. The EU AI Act and South Korea’s AI Framework Act demonstrate that transparency obligations and human oversight requirements can coexist with large-scale automated adjudication; the choice is not between efficiency and rights, but between efficiency with accountability and efficiency without it.

One limitation remains honestly on the table. That Congress can close this gap is not a promise that it will — immigration legislation is notoriously difficult to move, and the appetite to constrain the executive’s own tools may be scarce. But whether a reform is politically feasible and whether it is legally justified are distinct questions, and this paper defends the latter: what the law ought to require, not a forecast of when Congress will require it.

C. Toward a Rights-Based Framework for Algorithmic Adjudication

The deeper argument of this paper is not technical — it is normative. The adoption of AI tools in immigration adjudication has been framed, institutionally and in public discourse, as a question of administrative modernization: how can government agencies process more applications, more quickly, with fewer resources? That framing is not wrong, but it is incomplete. It treats efficiency as the primary value and human rights as a constraint to be managed, rather than recognizing that the legitimacy of an immigration system depends not only on its throughput but on its fairness.

The EU and South Korean frameworks reflect not merely that transparency and efficiency can coexist, but a deliberate ordering of priorities. They begin from the premise that AI systems affecting significant personal interests must be designed to protect rights first and optimize efficiency second. They embed that priority in legally enforceable obligations: explainability rights, audit mandates, human oversight requirements, and penalties for non-compliance.

The United States has, as of 2026, no equivalent framework for immigration AI. That absence is not a temporary gap waiting to be filled by future legislation — it is an active choice, reflected in the 2025 replacement of the prior AI-governance framework with one that creates no rights enforceable against the government, that imposes the costs of algorithmic error and discrimination on the population least equipped to bear them. The five recommendations offered here are not a comprehensive solution; they are a floor — the minimum legal infrastructure necessary to ensure that the efficiency gains of AI-assisted adjudication do not come at the cost of the rights of the people it processes.

This paper's argument stands or falls on its own terms, independent of its author's history. I note only that I write from a particular vantage point — having been, in 2025, subject to the kind of automated, unexplained processing this paper examines — and offer that not as authority, but as the reason the question first seemed worth asking.

An immigration system that cannot explain its decisions, that cannot demonstrate it is non-discriminatory, and that offers no meaningful recourse when its automated tools produce unjust results is not simply an inefficient system in need of reform. It is a system in which the rule of law has been partially displaced by algorithmic administration — and in which the human beings most affected by that displacement have been given no legal tools to push back. None of this is abstract. Behind every flagged file and every unexplained denial is a person who was never told what a machine decided about their life — and who deserved, at the very least, a system willing to explain itself. The technology is already here; the accountability is not. The only real question left is whether the law will catch up before the next two thousand people find themselves, as I once did, unable to learn why.

Notes

[1] Ashley Mowreader, “International Student Visas Revoked,” *Inside Higher Ed*, April 8, 2025 (updated April 25, 2025), <https://www.insidehighered.com/news/global/international-students-us/2025/04/07/where-students-have-had-their-visas-revoked>; Ashley Mowreader, “Federal Government Reverses SEVIS Terminations,” *Inside Higher Ed*, April 25, 2025, <https://www.insidehighered.com/news/global/international-students-us/2025/04/25/ice-reverses-course-sevis-terminations> (reporting “over 1,840 students and recent graduates from more than 280 colleges and universities” as of April 25, 2025).

[2] *Jane Doe 1 v. Bondi*, No. 1:25-cv-01998 (N.D. Ga. filed Apr. 11, 2025) (TRO entered Apr. 18, 2025, by Judge Victoria M. Calvert, ordering reinstatement of plaintiffs’ SEVIS records retroactive to March 31, 2025). See also ACLU of Georgia, “Temporary Restraining Order Granted in Student Visa Revocation Case *Doe v. Bondi*,” April 21, 2025, <https://www.acluga.org/press-releases/temporary-restraining-order-granted-student-visa-revocation-case-doe-v-bondi>.

[3] U.S. Department of Homeland Security, “AI Use Case Inventory,” January 28, 2026, <https://www.dhs.gov/ai/use-case-inventory/uscis>; Oleg Gherasimov, “USCIS Artificial Intelligence: What It’s Doing to Your Immigration Case Right Now,” SG Legal Group, April 23, 2026, <https://www.sglegalgroup.com/blog/uscis-artificial-intelligence-what-it-s-doing-to-your-immigration-case-right-now> (reporting “roughly 29 USCIS AI use cases” in the January 28, 2026, DHS inventory).

[4] The figures for the ELIS Evidence Classifier are reported in ACT-IAC, *National Use Case and Solutions Library*, “Evidence Classifier,” October 2022, <https://www.actiac.org/et-use-case/evidence-classifier> (estimating, for the period September 28, 2021–May 20, 2022, some 24 million page scrolls saved—approximately 13,348 hours—and an increase in the share of cases adjudicated within 30 days from about 30% to 58%), drawing on U.S. Citizenship and Immigration Services, *Privacy Impact Assessment for the USCIS Electronic Immigration System (USCIS ELIS)*, DHS/USCIS/PIA-056,

<https://www.dhs.gov/publication/dhsuscispia-056-uscis-electronic-immigration-system-uscis-elis>. These figures reflect processing efficiency, not adjudicative accuracy; USCIS has published no error-rate data for the classifier. The broader pattern of rising requests for evidence and denials is reported in Cozen O'Connor, "Growing Use of Artificial Intelligence in U.S. Immigration Adjudications Is Driving Higher RFE and Denial Rates," April 27, 2026, <https://www.cozen.com/news-resources/publications/2026/growing-use-of-artificial-intelligence-in-u-s-immigration-adjudications-is-driving-higher-rfe-and-denial-rates>.

[5] Caroline Haskins, "ICE Is Paying Palantir \$30 Million to Build 'ImmigrationOS' Surveillance Platform," *Wired*, April 18, 2025, <https://www.wired.com/story/ice-palantir-immigrationos>; American Civil Liberties Union, "All the Ways Palantir Is Assisting Trump's Abusive Removal Campaign," March 24, 2026, <https://www.aclu.org/news/privacy-technology/palantir-deportation-roundup>; American Immigration Council, "ICE to Use ImmigrationOS by Palantir, a New AI System, to Track Immigrants' Movements," August 21, 2025, <https://www.americanimmigrationcouncil.org/blog/ice-immigrationos-palantir-ai-track-immigrants>. Federal Contract ID: 7OCTD022FR0000170.

[6] Cozen O'Connor, "Growing Use of Artificial Intelligence."

[7] U.S. Citizenship and Immigration Services, "Form I-140 by Fiscal Year, Quarter and Case Status" (FY2025, Q1–Q3), <https://www.uscis.gov/tools/reports-and-studies/immigration-and-citizenship-data> (National Interest Waiver adjudicated-denial rate rising from 37.3% in Q1 to 46.0% in Q3 FY2025, and 38.8% in aggregate for Q1–Q3; denial rate computed as denials divided by adjudicated petitions). For the fourth quarter of FY2025, a National Foundation for American Policy analysis of subsequent USCIS data reports a 64.3% denial rate: Stuart Anderson, "U.S. Immigration Service Increases Denials for High-Skilled Immigrants," *Forbes*, April 22, 2026, <https://www.forbes.com/sites/stuartanderson/2026/04/22/us-immigration-service-increases-denials-for-high-skilled-immigrants/>. Available analyses attribute the increase principally to a stricter application of the *Dhanasar* standard; no published USCIS data isolates the contribution, if any, of AI-assisted screening to the trend.

[8] U.S. Citizenship and Immigration Services, “All USCIS Application and Petition Form Types (Fiscal Year 2025, Quarter 3),” Immigration and Citizenship Data, <https://www.uscis.gov/tools/reports-and-studies/immigration-and-citizenship-data> (reporting over 11 million pending cases across all form types as of mid-2025).

[9] U.S. Department of Homeland Security, “AI Use Case Inventory”; *Nextgov/FCW*, “Law Enforcement Is the Leading DHS Use Case for AI,” January 29, 2026, <https://www.nextgov.com/artificial-intelligence/2026/01/law-enforcement-leading-dhs-use-case-ai/411063> (reporting 238 total DHS AI use cases in the January 28, 2026, inventory).

[10] Gherasimov, “USCIS Artificial Intelligence.”

[11] ACT-IAC, “Evidence Classifier” (page-scroll and 30-day adjudication figures); Cozen O’Connor, “Growing Use of Artificial Intelligence” (classifier mis-tagging and resulting requests for evidence).

[12] BMD LLC, “Invisible Algorithms: The Hidden Role of Artificial Intelligence in USCIS Immigration Processing,” February 10, 2026, <https://www.bmdllc.com/resources/blog/invisible-algorithms-the-hidden-role-of-artificial-intelligence-in-uscis-immigration-processing>.

[13] Haskins, “ICE Is Paying Palantir”; ACLU, “All the Ways Palantir Is Assisting.”

[14] Greenspoon Marder LLP, “How Artificial Intelligence Is Revolutionizing Immigration Compliance and Enforcement in 2026,” January 7, 2026, <https://www.gmlaw.com/news/how-artificial-intelligence-is-revolutionizing-immigration-compliance-and-enforcement-in-2026>.

[15] Greenberg Traurig LLP, “The U.S. Government’s Expanding Use of AI in Immigration Adjudications: Considerations for Employers and Investors,” December 18, 2025, <https://www.gtlaw-insidebusinessimmigration.com/emerging-technology/the-us-governments-expanding-use-of-ai-in-immigration-adjudications-considerations-for-employers-and-investors>.

[16] Greenspoon Marder, “How Artificial Intelligence Is Revolutionizing Immigration Compliance.”

[17] Gherasimov, “USCIS Artificial Intelligence.”

[18] Cozen O’Connor, “Growing Use of Artificial Intelligence.”

[19] U.S. Citizenship and Immigration Services, “Form I-140 by Fiscal Year, Quarter and Case Status” (FY2025), <https://www.uscis.gov/tools/reports-and-studies/immigration-and-citizenship-data> (NIW adjudicated-denial rate rising from 37.3% in Q1 to 46.0% in Q3 FY2025); for the fourth quarter of FY2025, see the National Foundation for American Policy analysis reported in Stuart Anderson, “U.S. Immigration Service Increases Denials for High-Skilled Immigrants,” *Forbes*, April 22, 2026, <https://www.forbes.com/sites/stuartanderson/2026/04/22/us-immigration-service-increases-denials-for-high-skilled-immigrants/> (64.3% NIW denial rate; EB-1A rising from 25.6% to 46.6%). On the attribution of the increase to a more exacting application of the *Dhanasar* standard rather than to automation, see *id.* (remarks of Efren Hernandez, former USCIS adjudicator). The operational mechanisms described—mis-tagged evidence and cross-document mismatch flags—are reported in Cozen O’Connor, “Growing Use of Artificial Intelligence.”

[20] Richard T. Herman, “Algorithmic Immigration: Is Artificial Intelligence Shaping USCIS Decisions in 2026?,” Herman Legal Group, accessed August 4, 2026, <https://www.lawfirm4immigrants.com/uscis-artificial-intelligence-2026>

[21] *Kleindienst v. Mandel*, 408 U.S. 753, 766 (1972).

[22] *Mathews v. Diaz*, 426 U.S. 67, 77 (1976); *Plyler v. Doe*, 457 U.S. 202, 215 (1982); *Zadvydas v. Davis*, 533 U.S. 678, 693 (2001).

[23] Under the APA, judicial review is available only for “final agency action for which there is no other adequate remedy in a court.” 5 U.S.C. § 704. The Supreme Court has established a two-part test for finality: the action must mark the consummation of the agency’s decision-making process, and it must be one by which rights or

obligations have been determined or from which legal consequences will flow. *Bennett v. Spear*, 520 U.S. 154, 177–78 (1997); see also *Lujan v. National Wildlife Federation*, 497 U.S. 871, 882 (1990). Because AI systems in immigration adjudication are formally characterized as decision-support tools rather than decision-makers, their outputs are unlikely to satisfy the *Bennett* finality test, insulating algorithmic processing from direct APA review.

[24] David A. Martin, “Why Immigration’s Plenary Power Doctrine Endures,” *Oklahoma Law Review* 68 (2015): 1.

[25] Regulation (EU) 2016/679 (GDPR), art. 22.

[26] Case C-634/21, *SCHUFA Holding AG*, ECLI:EU:C:2023:957 (ECJ, Dec. 7, 2023).

[27] Regulation (EU) 2024/1689 (EU AI Act), art. 86.

[28] EU AI Act, art. 86; see also “Article 86—Right to Explanation of Individual Decision-Making,” artificialintelligenceact.eu, <https://artificialintelligenceact.eu/article/86> (annotated text of the provision); European Parliament, “EU Artificial Intelligence Act: First Regulation on Artificial Intelligence,” June 2023, <https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence> (noting geographic scope limitations and phased implementation); the high-risk regime’s application to standalone Annex III systems was subsequently deferred from August 2026 to December 2027 by the Digital Omnibus, Regulation (EU) 2026/1744.

[29] Framework Act on Artificial Intelligence Development and Establishment of a Foundation for Trustworthiness (AI Framework Act), Law No. 20676 (Jan. 21, 2025) (S. Kor.); Shin & Kim, “AI Basic Act Update: Enforcement and Key Implications,” February 11, 2026, <https://www.shinkim.com/eng/media/newsletter/3117>.

[30] AI Framework Act, art. 2(4).

[31] *Id.* art. 31(1).

[32] *Id.* arts. 32, 34.

[33] Juho Yoon, Jiyoung Sohn, and Jeonghee Kang, “Significance of the Passage of the AI Framework Act and Its Impact on the Industry,” Bae, Kim & Lee LLC, Lexology, January 10, 2025, <https://www.lexology.com/library/detail.aspx?g=9f68b5ab-b6b8-4822-823a-4872a8152cfb>.

[34] Natalie Sheard, “Algorithm-Facilitated Discrimination: A Socio-Legal Study of the Use by Employers of Artificial Intelligence Hiring Systems,” *Journal of Law and Society* 52, no. 2 (2025): 269–291, <https://doi.org/10.1111/jols.12535>.

[35] Yoann Emian, “Digital Borders and Racial Codes in AI Migration Control,” Ontario Bar Association, July 2, 2025, <https://oba.org/digital-borders-and-racial-codes-in-ai-migration-control>

[36] Letter from Electronic Frontier Foundation, Just Futures Law, The Advocates for Human Rights, and 140+ Organizations to Secretary Alejandro Mayorkas, “Cancel DHS Use of AI Technologies for Immigration Enforcement and Adjudication by December 1, 2024,” September 4, 2024, https://fedscoop.com/wp-content/uploads/sites/5/2024/09/Letter-from-140-Orgs-to-DHS_-Cancel-Use-of-AI-by-Dec-1.pdf; see also EFF, “EFF and 140 Other Organizations Call for an End to AI Use in Immigration Decisions,” September 5, 2024, <https://www.eff.org/deeplinks/2024/09/eff-140-other-organizations-call-end-ai-use-immigration-decisions>; Rebecca Heilweil, “140+ Groups Flag DHS’s AI Use Cases in New Letter to Mayorkas,” *FedScoop*, September 4, 2024, <https://fedscoop.com/dhs-ai-use-cases-letter-mayorkas>.

[37] Rachel Levinson-Waldman and Spencer Reynolds, “A Start for AI Transparency at DHS with Room to Grow,” Brennan Center for Justice, January 22, 2025, <https://www.brennancenter.org/our-work/analysis-opinion/start-ai-transparency-dhs-room-grow> (assessing DHS’s AI use-case inventory and finding its disclosures about the ICE Hurricane Score incomplete); Just Futures Law and Mijente, *Automating Deportation: The Artificial Intelligence Behind the Department of Homeland Security’s Immigration Enforcement Regime* (June 2024),

<https://mijente.net/wp-content/uploads/2024/06/Automating-Deportation.pdf> (documenting that DHS deploys these tools without notifying impacted individuals).

[38] International Covenant on Civil and Political Rights, art. 26, Dec. 16, 1966, 999 U.N.T.S. 171. The United States ratified the Covenant in 1992 subject to a declaration that its substantive provisions (Articles 1 through 27) are not self-executing. 138 Cong. Rec. S4781-84 (daily ed. Apr. 2, 1992); see *Sosa v. Alvarez-Machain*, 542 U.S. 692, 728 (2004) (noting the declaration).

[39] Zihan Ma, "The Inadequacy of the Current International Human Rights Regime for Algorithm Discrimination," *Michigan Journal of International Law* (blog), February 2024, <https://www.mjilonline.org/the-inadequacy-of-the-current-international-human-rights-regime-for-algorithm-discrimination>.

[40] Henry J. Friendly, "Some Kind of Hearing," *University of Pennsylvania Law Review* 123 (1975): 1267; see also *Pangea Legal Services v. USCIS*, No. 1:24-cv-02809 (D.D.C. filed Oct. 3, 2024) (FOIA action seeking disclosure of DHS and USCIS AI use policies and system documentation), <https://www.justfutureslaw.org/legal-filings/dhsaifolia>.

[41] Pieter Binst et al., "Automated Decision-Making and Artificial Intelligence at European Borders and Their Risks for Human Rights," arXiv preprint arXiv:2410.17278 (2024), <https://arxiv.org/abs/2410.17278>.

[42] *Houston Federation of Teachers v. Houston Independent School District*, 251 F. Supp. 3d 1168, 1177 (S.D. Tex. 2017).

[43] Exec. Order No. 14179, 90 Fed. Reg. 8741 (Jan. 23, 2025) (directing OMB to revise Memoranda M-24-10 and M-24-18 within sixty days). Pursuant to that order, OMB Memorandum M-25-21, "Accelerating Federal Use of AI through Innovation, Governance, and Public Trust" (Apr. 3, 2025), rescinded and replaced M-24-10, consolidating its separate "rights-impacting" and "safety-impacting" categories into a single "high-impact AI" definition. See M-25-21 § 6(k)

(treating as presumptively high-impact AI used in risk assessments relating to a foreign national’s entry into or continued presence in the United States, including immigration, asylum, detention, and travel-authorization determinations); *id.* § 4(b)(vi) (instructing agencies to offer consistent remedies or appeals); *id.* n.4 (providing that the memorandum applies only to agencies’ own use of AI and creates no right or benefit, substantive or procedural, enforceable against the United States). The use-case inventory obligation survives the transition: *id.* § 3(b)(v) (requiring each agency to update and publicly post its AI use-case inventory at least annually), implementing the Advancing American AI Act, 40 U.S.C. § 7225.

[44] BMD LLC, “Invisible Algorithms.”

[45] EU AI Act, art. 86.

[46] Binst et al., “Automated Decision-Making.”

[47] Cozen O’Connor, “Growing Use of Artificial Intelligence.”

[48] White House Office of Science and Technology Policy, *Blueprint for an AI Bill of Rights: Making Automated Systems Work for the American People* (October 2022), <https://bidenwhitehouse.archives.gov/ostp/ai-bill-of-rights>.

[49] *Houston Federation of Teachers*, 251 F. Supp. 3d at 1177.

[50] Littler, “What Does the 2025 Artificial Intelligence Legislative and Regulatory Landscape Look Like for Employers?,” January 14, 2025, <https://www.littler.com/news-analysis/asap/what-does-2025-artificial-intelligence-legislative-and-regulatory-landscape-look>.

[51] GDPR, art. 22.

[52] U.S. Equal Employment Opportunity Commission, “Select Issues: Assessing Adverse Impact in Software, Algorithms, and Artificial Intelligence Used in Employment Selection Procedures Under Title VII of the Civil Rights Act of 1964,” Technical Assistance Document EEOC-NVTA-2023-2, May 18, 2023, <https://www.eeoc.gov/laws/guidance/questions-and-answers-clarify-and-provide-a-common-interpretation-uniform-guidelines>.

Bibliography

ACLU of Georgia. “Temporary Restraining Order Granted in Student Visa Revocation Case *Doe v. Bondi*.” April 21, 2025.

<https://www.acluga.org/press-releases/temporary-restraining-order-granted-student-visa-revocation-case-doe-v-bondi>.

ACT-IAC. “Evidence Classifier.” *National Use Case and Solutions Library*, October 2022. <https://www.actiac.org/et-use-case/evidence-classifier>.

American Civil Liberties Union. “All the Ways Palantir Is Assisting Trump’s Abusive Removal Campaign.” March 24, 2026. <https://www.aclu.org/news/privacy-technology/palantir-deportation-roundup>.

American Immigration Council. “ICE to Use ImmigrationOS by Palantir, a New AI System, to Track Immigrants’ Movements.” August 21, 2025. <https://www.americanimmigrationcouncil.org/blog/ice-immigrationos-palantir-ai-track-immigrants>.

Anderson, Stuart. “U.S. Immigration Service Increases Denials for High-Skilled Immigrants.” *Forbes*, April 22, 2026. <https://www.forbes.com/sites/stuartanderson/2026/04/22/us-immigration-service-increases-denials-for-high-skilled-immigrants/>.

“Article 86—Right to Explanation of Individual Decision-Making.” *artificialintelligenceact.eu*. <https://artificialintelligenceact.eu/article/86>.

Binst, Pieter, et al. “Automated Decision-Making and Artificial Intelligence at European Borders and Their Risks for Human Rights.” arXiv preprint arXiv:2410.17278 (2024). <https://arxiv.org/abs/2410.17278>.

BMD LLC. “Invisible Algorithms: The Hidden Role of Artificial Intelligence in USCIS Immigration Processing.” February 10, 2026. <https://www.bmdllc.com/resources/blog/invisible-algorithms-the-hidden-role-of-artificial-intelligence-in-uscis-immigration-processing>.

Cozen O’Connor. “Growing Use of Artificial Intelligence in U.S. Immigration Adjudications Is Driving Higher RFE and Denial Rates.” April 27, 2026. <https://www.cozen.com/news-resources/publications/2026/growing-use-of-artificial-intelligence-in-u-s-immigration-adjudications-is-driving-higher-rfe-and-denial-rates>.

Electronic Frontier Foundation. “EFF and 140 Other Organizations Call for an End to AI Use in Immigration Decisions.” September 5, 2024. <https://www.eff.org/deeplinks/2024/09/eff-140-other-organizations-call-end-ai-use-immigration-decisions>.

Electronic Frontier Foundation, Just Futures Law, The Advocates for Human Rights, and 140+ Organizations. Letter to Secretary Alejandro Mayorkas, “Cancel DHS Use of AI Technologies for Immigration Enforcement and Adjudication by December 1, 2024.” September 4, 2024. https://fedscoop.com/wp-content/uploads/sites/5/2024/09/Letter-from-140-Orgs-to-DHS_-Cancel-Use-of-AI-by-Dec-1.pdf.

Emian, Yoann. “Digital Borders and Racial Codes in AI Migration Control.” Ontario Bar Association, July 2, 2025. <https://oba.org/digital-borders-and-racial-codes-in-ai-migration-control>.

European Parliament. “EU Artificial Intelligence Act: First Regulation on Artificial Intelligence.” June 2023. <https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence>.

Friendly, Henry J. “Some Kind of Hearing.” *University of Pennsylvania Law Review* 123 (1975): 1267.

Gherasimov, Oleg. “USCIS Artificial Intelligence: What It’s Doing to Your Immigration Case Right Now.” SG Legal Group, April 23, 2026. <https://www.sglegalgroup.com/blog/uscis-artificial-intelligence-what-its-doing-to-your-immigration-case-right-now>.

Greenspoon Marder LLP. “How Artificial Intelligence Is Revolutionizing Immigration Compliance and Enforcement in 2026.” January 7, 2026. <https://www.gmlaw.com/news/how-artificial-intelligence-is-revolutionizing-immigration-compliance-and-enforcement-in-2026>.

Haskins, Caroline. “ICE Is Paying Palantir \$30 Million to Build ‘ImmigrationOS’ Surveillance Platform.” *Wired*, April 18, 2025. <https://www.wired.com/story/ice-palantir-immigrationos>.

Heilweil, Rebecca. “140+ Groups Flag DHS’s AI Use Cases in New Letter to Mayorkas.” *FedScoop*, September 4, 2024. <https://fedscoop.com/dhs-ai-use-cases-letter-mayorkas>.

Herman, Richard T. “Algorithmic Immigration: Is Artificial Intelligence Shaping USCIS Decisions in 2026?” Herman Legal Group. Accessed August 4, 2026. <https://www.lawfirm4immigrants.com/uscis-artificial-intelligence-2026>.

Just Futures Law and Mijente. *Automating Deportation: The Artificial Intelligence Behind the Department of Homeland Security’s Immigration Enforcement Regime*. June 2024. <https://mijente.net/wp-content/uploads/2024/06/Automating-Deportation.pdf>.

- Kalmykov, Kate. “The U.S. Government’s Expanding Use of AI in Immigration Adjudications: Considerations for Employers and Investors.” Greenberg Traurig LLP, December 18, 2025. <https://www.gtlaw-insidebusinessimmigration.com/emerging-technology/the-us-governments-expanding-use-of-ai-in-immigration-adjudications-considerations-for-employers-and-investors>.
- Levinson-Waldman, Rachel, and Spencer Reynolds. “A Start for AI Transparency at DHS with Room to Grow.” Brennan Center for Justice, January 22, 2025. <https://www.brennancenter.org/our-work/analysis-opinion/start-ai-transparency-dhs-room-grow>.
- Littler. “What Does the 2025 Artificial Intelligence Legislative and Regulatory Landscape Look Like for Employers?” January 14, 2025. <https://www.littler.com/news-analysis/asap/what-does-2025-artificial-intelligence-legislative-and-regulatory-landscape-look>.
- Ma, Zihan. “The Inadequacy of the Current International Human Rights Regime for Algorithm Discrimination.” *Michigan Journal of International Law* (blog), February 2024. <https://www.mjilonline.org/the-inadequacy-of-the-current-international-human-rights-regime-for-algorithm-discrimination>.
- Martin, David A. “Why Immigration’s Plenary Power Doctrine Endures.” *Oklahoma Law Review* 68 (2015): 1.
- Mowreader, Ashley. “Federal Government Reverses SEVIS Terminations.” *Inside Higher Ed*, April 25, 2025. <https://www.insidehighered.com/news/global/international-students-us/2025/04/25/ice-reverses-course-sevis-terminations>.

Mowreader, Ashley. “International Student Visas Revoked.” *Inside Higher Ed*, April 8, 2025 (updated April 25, 2025). <https://www.insidehighered.com/news/global/international-students-us/2025/04/07/where-students-have-had-their-visas-revoked>.

Nextgov/FCW. “Law Enforcement Is the Leading DHS Use Case for AI.” <https://www.eeoc.gov/laws/guidance/questions-and-answers-clarify-and-provide-a-common-interpretation-uniform-guidelines>.

Sheard, Natalie. “Algorithm-Facilitated Discrimination: A Socio-Legal Study of the Use by Employers of Artificial Intelligence Hiring Systems.” *Journal of Law and Society* 52, no. 2 (2025): 269–91. <https://doi.org/10.1111/jols.12535>.

Shin & Kim. “AI Basic Act Update: Enforcement and Key Implications.” February 11, 2026. <https://www.shinkim.com/eng/media/newsletter/3117>.

U.S. Citizenship and Immigration Services. “All USCIS Application and Petition Form Types (Fiscal Year 2025, Quarter 3).” Immigration and Citizenship Data. <https://www.uscis.gov/tools/reports-and-studies/immigration-and-citizenship-data>.

U.S. Citizenship and Immigration Services. “Form I-140 by Fiscal Year, Quarter and Case Status.” FY2025. <https://www.uscis.gov/tools/reports-and-studies/immigration-and-citizenship-data>.

U.S. Citizenship and Immigration Services. *Privacy Impact Assessment for the USCIS Electronic Immigration System (USCIS ELIS)*. DHS/USCIS/PIA-056. <https://www.dhs.gov/publication/dhsuscispia-056-uscis-electronic-immigration-system-uscis-elis>.

U.S. Department of Homeland Security. “AI Use Case Inventory.” January 28, 2026. <https://www.dhs.gov/ai/use-case-inventory/uscis>.

U.S. Equal Employment Opportunity Commission. “Select Issues: Assessing Adverse Impact in Software, Algorithms, and Artificial Intelligence Used in Employment Selection Procedures Under Title VII of the Civil Rights Act of 1964.” Technical Assistance Document EEOC-NVTA-2023-2, May 18, 2023.

White House Office of Science and Technology Policy. *Blueprint for an AI Bill of Rights: Making Automated Systems Work for the American People*. October 2022. <https://bidenwhitehouse.archives.gov/ostp/ai-bill-of-rights>.

Yoon, Juho, Jiyoung Sohn, and Jeonghee Kang. “Significance of the Passage of the AI Framework Act and Its Impact on the Industry.” Bae, Kim & Lee LLC, Lexology, January 10, 2025. <https://www.lexology.com/library/detail.aspx?g=9f68b5ab-b6b8-4822-823a-4872a8152cfb>.